

EXHIBIT 18

TPU locations

Accelerators, such as Tensor Processing Units (TPUs), are specialized hardware components that can significantly speed up machine learning and high performance computing workloads. When you create Compute Engine instances or other resources on Google Cloud, you can attach accelerators to them. For an overview of these accelerators, see [TPU machines](/compute/docs/tpus/tpu-machines) (/compute/docs/tpus/tpu-machines).

The availability of specific TPU models varies by region and zone. This document provides a comprehensive list of accelerator availability to help you plan your deployments and select the appropriate locations for your resources.

The table in this document focuses on TPU availability for Compute Engine. If you use TPUs with other Google Cloud services, accelerator availability might differ. For those services, see the following documentation:

- [Google Kubernetes Engine \(GKE\): Validate TPU availability in GKE](/kubernetes-engine/docs/concepts/plan-tpus#availability) (/kubernetes-engine/docs/concepts/plan-tpus#availability)
- [Vertex AI: accelerator locations](/vertex-ai/docs/general/locations#accelerators) (/vertex-ai/docs/general/locations#accelerators)

Zones that support TPUs

Google Cloud offers TPU resources in multiple locations worldwide. These locations consist of the standard regions and zones, which you can read about on the [Regions and zones](/compute/docs/regions-zones) (/compute/docs/regions-zones) page. Each zone offers a variety of TPU versions. When you create resources that use accelerators—such as Compute Engine instances or GKE node pools—you must select a zone where that accelerator type is available. Compute Engine supports TPU7x, TPU v6e, and TPU v5p.

AI zones

Important: Access to AI zones is restricted unless you've enabled them for your Google Cloud project. To use them, [enable AI zones](/compute/docs/regions-zones/manage-ai-zones#enable_ai_zones) (/compute/docs/regions-zones/manage-ai-zones#enable_ai_zones) ([Preview](https://cloud.google.com/products#product-launch-stages) (https://cloud.google.com/products#product-launch-stages)).

In addition to standard zones, TPUs are available in AI zones. These Google Cloud locations are specifically designed to house resources for large-scale AI and machine learning workloads. AI zones offer access to TPUs, and might also offer other non-accelerator machine types to support the primary AI and machine learning workloads that run on the accelerators. To learn more about these specialized zones, see [About AI zones](#) (/compute/docs/regions-zones/ai-zones). To view accelerator availability by zone type, use the **Select a zone type** filter.

TPU availability

The following sortable table lets you select different options to see where TPU resources are available on Compute Engine. You can search by location, or TPU version. Note that Compute Engine supports TPU7x, TPU v6e, and TPU v5p.

Select a location ▾

Select a zone type ▾

Select a TPU version ▾

Clear all

Zone	Location	Zone type	TPU version
europa-west4-b	Eemshaven, Netherlands, Europe	Standard	• v5p
us-central1-a	Council Bluffs, Iowa, North America	Standard	• v5p
us-east5-a	Columbus, Ohio, North America	Standard	• v6e (Trillium) • v5p

Feature availability by region and zone

In addition to the availability of specific accelerator models, some related features and configurations have their own regional or zonal restrictions. The following sections outline the availability for these specific features.

Future reservations in calendar mode and Flex-start VMs

Future reservations in calendar mode and TPU Flex-start VMs are only available in specific zones.

- For a list of zones that support calendar mode reservations for TPUs, see [TPU limitations for calendar mode reservations](#) (/compute/docs/instances/future-reservations-calendar-mode-overview#limitations-tpu).
- For a list of zones that support TPU Flex-start VMs, see [Limitations for TPU Flex-start VMs](#) (/compute/docs/instances/about-flex-start-vm#limitations-for-machine-types).

What's next?

- [Learn more about geography and regions](#) (/docs/geography-and-regions).
- [Learn about TPU pricing](#) (/tpu/pricing).
- [Learn about TPU resources in Compute Engine](#) (/compute/docs/tpus/tpu-resources-in-compute-engine).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (https://creativecommons.org/licenses/by/4.0/), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (https://www.apache.org/licenses/LICENSE-2.0). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (https://developers.google.com/site-policies). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2026-07-28 UTC.